

AI Trust & Security Readiness for RAG and Autonomous Agents

Evidence-based readiness, runtime controls, and launch-gate decisions for AI systems that retrieve, reason, and act.

RAG Readiness	Agent Readiness	Launch Gates	Client Evidence Template
1. Problem AI systems often move from prototype to decision pressure before evidence is complete. This creates risk across retrieval scope, identity controls, tool usage, and answer supportability.		2. Approach Use an evidence-based readiness methodology: define controls, run validation, collect artifacts, map residual risk, and enforce launch gates.	
3. Evidence Chain Local runtime evidence harnesses; CI artifact proof; observability correlation evidence; local agent runtime evidence; a staging-demo evidence model; and a client-specific production template for real evidence intake.		4. What This Proves This package demonstrates a structured, auditable readiness methodology and launch-gate decision model for AI trust and security planning.	
5. What This Does Not Prove This package does not prove production readiness, client approval, compliance approval, or real production GO.		6. Client Next Step Run client discovery, map identity/policy boundaries, and populate Version 4 with real artifacts, runtime proof, residual-risk ownership, and named approvals.	
7. Production Boundary		Production readiness remains blocked until real client data, users, policies, runtime evidence, compliance review, residual-risk acceptance, and named approval are verified.	
Next client action: discovery call -> evidence intake -> Version 4 client template -> launch-gate review.			

Portfolio/readiness methodology - not production-ready without real client evidence. Production readiness and GO decision remain blocked until real client evidence is verified.